

Sample RAG Document

This is a sample PDF created for testing a PDF loader in a Retrieval-Augmented Generation project.

Artificial intelligence allows software systems to understand text, answer questions, and retrieve relevant information from large collections of documents.

A typical RAG pipeline contains several steps: document loading, text splitting, embedding generation, vector database storage, retrieval, and response generation.

FastAPI can be used to expose the RAG system through REST APIs, while PostgreSQL can store application data such as users, documents, and conversations.